

Bounding the “AGI” Debate:

Why forecasting its impact is a hard task.

Filotas Theodosiou

Different definitions, different forecasts.

- Many mentions of “AGI”.
- (Almost) all were different.
- Each favoured the topic at hand.

International Conference on Large-Scale AI Risks

📅 May 26, 2025 12:30 PM to May 28, 2025 01:15 PM (Europe/Brussels / UTC200)

📍 Institute of Philosophy

👤 Lode Lauwaert

📅 iCal

AI's Large-Scale Risks: Control, Governance, and Ethics

While advanced AI has large beneficial potential that attracts billions in both private and public funding, the academic community interested in the understanding and mitigation of the technology's large-scale risks, such as enabling bioterrorism and the loss of human control over the technology, is still tiny in comparison. Addressing this important global problem requires experts from various disciplines to cooperate on the identification of issues and solutions. While the academic community has recently experienced growth, it is still largely centered around the US and UK. Moreover, disciplines remain largely disconnected, since they present their work at their specific venues. This academic conference on AI's large-scale risks brings together experts from various disciplines to discuss technical, social, and philosophical topics around the control, governance, and ethics of advanced AI. Putting a strong emphasis on the interdisciplinary nature of the topic, the conference will include an interdisciplinary panel discussion on potential symbiotic effects of interdisciplinary AI risk research.

Different definitions, different forecasts.

Different definitions of AGI



OpenAI

“outperform humans at most economically valuable work”



DeepMind

“as capable as humans at most cognitive tasks”



AWS

“human-like intelligence” and “self-teach”



ARC-AGI

“match the learning efficiency of humans”



IBM

“match or exceed human cognitive abilities across any task”



Anthropic

“transformative AI” and “human-level performance at most intellectual tasks”

Different definitions, different forecasts.

Different definitions of AGI



OpenAI

“outperform humans at most economically valuable work”



DeepMind

“as capable as humans at most cognitive tasks”



AWS

“human-like intelligence” and “self-teach”



ARC-AGI

“match the learning efficiency of humans”



IBM

“match or exceed human cognitive abilities across any task”



Anthropic

“transformative AI” and “human-level performance at most intellectual tasks”

Different types of forecasts serve different discussions

AGI as economic work



forecast jobs and productivity

AGI as cognitive breadth



forecast capability timelines

AGI as transformative AI



forecast societal disruption


Different definitions, different forecasts.

Forecasting the Impact of Artificial General Intelligence (AGI)


SPYROS MAKRIDAKIS

PREVIEW *It's now more than seven years since Foresight published a five-part series by Spyros Makridakis on the impact of Artificial Intelligence. In those articles, he repeatedly warned against linear thinking and the serious underestimation of exponential technological progress. Yet, like nearly everyone else, Spyros himself was far off the mark regarding both the speed and scope of AI's advances. What he described in 2018 as "plausibly decades away" has arrived in prototype form in 2025. Frontier models now routinely demonstrate superior capabilities across language, coding, mathematics, science, medicine, and multimodal reasoning, and do so at and increasingly beyond the aptitude of most flesh-and-blood professionals. We stand at the threshold of Artificial General Intelligence (AGI): systems that approach the ability to perform many intellectual tasks that a human can accomplish, and do so at superhuman speed and scale. In this follow-up installment, Makridakis examines what the imminent arrival of AGI means for management and, by extension, for the profession of forecasting itself.*

What do people use to solve their intellectual tasks?






What do people use to solve their intellectual tasks?




What “AGI” will feel like for most people

What do people use to solve their intellectual tasks?






What “AGI” will feel like for most people

What do people use to solve their intellectual tasks?




Interact with an AI via a platform to solve a task.



What “AGI” will feel like for most people

How did we get here?



1

Machines

Automated
productivity tasks



2

Computers

Moved many tasks
into the digital world



3

Internet


Connected digital
systems



4

Cloud & Mobile

Made software
always available



5

AI Agents

Automate task execution
in the digital world

The general framework I am going to use.


Physical world




Human goals:

- Improve quality of life
- Earn money
- Publish more
- Other...

The general framework I am going to use.



The general framework I am going to use.



The general framework I am going to use.

We are still far from robust general embodied intelligence

- Data scarcity
- Real world is unforgiving
- Current models generalize poorly in open physical environments
- Hardware
- Safety/trust

Physical world



Human goals:


- Improve quality of life
- Earn money
- Publish more
- Other...

Instructions /
Feedback



Actions /
Clarification

Digital world



Digital tasks to achieve goals

-  Research
-  Content creation
-  Coding
-  Computer tasks

AGI
happens
here

Discussion points based on this framework

Forecasting the Impact of Artificial General Intelligence (AGI)

SPYROS MAKRIDAKIS

PREVIEW *It's now more than seven years since Foresight published a five-part series by Spyros Makridakis on the impact of Artificial Intelligence. In those articles, he repeatedly warned against linear thinking and the serious underestimation of exponential technological progress. Yet, like nearly everyone else, Spyros himself was far off the mark regarding both the speed and scope of AI's advances. What he described in 2018 as "plausibly decades away" has arrived in prototype form in 2025. Frontier models now routinely demonstrate superior capabilities across language, coding, mathematics, science, medicine, and multimodal reasoning, and do so at and increasingly beyond the aptitude of most flesh-and-blood professionals. We stand at the threshold of Artificial General Intelligence (AGI): systems that approach the ability to perform many intellectual tasks that a human can accomplish, and do so at superhuman speed and scale. In this follow-up installment, Makridakis examines what the imminent arrival of AGI means for management and, by extension, for the profession of forecasting itself.*

The AGI race has begun. We are building machines that can think and reason. By 2025/26, these machines will outpace many college graduates. By the end of the decade, they will be smarter than you or I; we will have superintelligence, in the true sense of the word.

Leopold Aschenbrenner (2024)

Artificial Intelligence will profoundly reshape modern society, transforming workplaces, entire industries, and our everyday lives while captivating – and in many cases unnerving – the world with its rapid advancements. Looking forward, AI's trajectory points to continued exponential growth, with many experts anticipating the emergence of Artificial General Intelligence (AGI) in the next years.

This paper aims to forecast the profound impact of AGI, serving as a sequel to my five-part *Foresight* series on forecasting the impact of AI (Makridakis, 2017; 2018a; 2018b; 2018c; 2018d). To illustrate the remarkable capabilities of contemporary AI, I provided a generative Large Language Model (LLM) the following prompt:

"I am Spyros Makridakis, and I want to publish a sequel to my five-part series on 'Forecasting the Impact of AI.' Could you

provide a short summary of the five published articles and suggest an outline for the sequel paper, 'Forecasting the Impact of AGI'?"

The LLM's response, reproduced in the Appendix that follows, was generated in mere seconds. It delivered concise, accurate summaries of each article, demonstrated a deep understanding of their content and interconnections, and proposed a structured outline for the new paper. Accomplishing this otherwise would have required a human expert spending considerable time, effort, and careful thinking.

THE ACCELERATING PATH OF AGI

With the benefit of hindsight, how did my 2017-18 predictions about the impact of AI hold up in February 2026? First, what I got remarkably right is the prediction that exponential progress would continue

Discussion points based on this framework

1. Why forecasting the impact of “AGI” is hard and why we should focus on what we have:
 - From capability to impact.
 - Three overlapping circles and a “naive forecast”.
 - Revenue as a predictor
 - Ownership and inequality.

Forecasting the Impact of Artificial General Intelligence (AGI)

SPYROS MAKRIDAKIS

PREVIEW It's now more than seven years since Foresight published a five-part series by Spyros Makridakis on the impact of Artificial Intelligence. In those articles, he repeatedly warned against linear thinking and the serious underestimation of exponential technological progress. Yet, like nearly everyone else, Spyros himself was far off the mark regarding both the speed and scope of AI's advances. What he described in 2018 as “plausibly decades away” has arrived in prototype form in 2025. Frontier models now routinely demonstrate superior capabilities across language, coding, mathematics, science, medicine, and multimodal reasoning, and do so at and increasingly beyond the aptitude of most flesh-and-blood professionals. We stand at the threshold of Artificial General Intelligence (AGI): systems that approach the ability to perform many intellectual tasks that a human can accomplish, and do so at superhuman speed and scale. In this follow-up installment, Makridakis examines what the imminent arrival of AGI means for management and, by extension, for the profession of forecasting itself.

The AGI race has begun. We are building machines that can think and reason. By 2025/26, these machines will outpace many college graduates. By the end of the decade, they will be smarter than you or I; we will have superintelligence, in the true sense of the word.

Leopold Aschenbrenner (2024)

Artificial Intelligence will profoundly reshape modern society, transforming workplaces, entire industries, and our everyday lives while captivating – and in many cases unnerving – the world with its rapid advancements. Looking forward, AI's trajectory points to continued exponential growth, with many experts anticipating the emergence of Artificial General Intelligence (AGI) in the next years.

This paper aims to forecast the profound impact of AGI, serving as a sequel to my five-part Foresight series on forecasting the impact of AI (Makridakis, 2017; 2018a; 2018b; 2018c; 2018d). To illustrate the remarkable capabilities of contemporary AI, I provided a generative Large Language Model (LLM) the following prompt:

‘I am Spyros Makridakis, and I want to publish a sequel to my five-part series on ‘Forecasting the Impact of AI.’ Could you

provide a short summary of the five published articles and suggest an outline for the sequel paper, ‘Forecasting the Impact of AGI?’

The LLM's response, reproduced in the Appendix that follows, was generated in mere seconds. It delivered concise, accurate summaries of each article, demonstrated a deep understanding of their content and interconnections, and proposed a structured outline for the new paper. Accomplishing this otherwise would have required a human expert spending considerable time, effort, and careful thinking.

THE ACCELERATING PATH OF AGI

With the benefit of hindsight, how did my 2017-18 predictions about the impact of AI hold up in February 2026? First, what I got remarkably right is the prediction that exponential progress would continue

Discussion points based on this framework

1. Why forecasting the impact of “AGI” is hard and why we should focus on what we have:
 - From capability to impact.
 - Three overlapping circles and a “naive forecast”.
 - Revenue as a predictor
 - Ownership and inequality.
2. What this framework means for forecasting

Forecasting the Impact of Artificial General Intelligence (AGI)

SPYROS MAKRIDAKIS

PREVIEW It's now more than seven years since Foresight published a five-part series by Spyros Makridakis on the impact of Artificial Intelligence. In those articles, he repeatedly warned against linear thinking and the serious underestimation of exponential technological progress. Yet, like nearly everyone else, Spyros himself was far off the mark regarding both the speed and scope of AI's advances. What he described in 2018 as “plausibly decades away” has arrived in prototype form in 2025. Frontier models now routinely demonstrate superior capabilities across language, coding, mathematics, science, medicine, and multimodal reasoning, and do so at and increasingly beyond the aptitude of most flesh-and-blood professionals. We stand at the threshold of Artificial General Intelligence (AGI): systems that approach the ability to perform many intellectual tasks that a human can accomplish, and do so at superhuman speed and scale. In this follow-up installment, Makridakis examines what the imminent arrival of AGI means for management and, by extension, for the profession of forecasting itself.

The AGI race has begun. We are building machines that can think and reason. By 2025/26, these machines will outpace many college graduates. By the end of the decade, they will be smarter than you or I; we will have superintelligence, in the true sense of the word.

Leopold Aschenbrenner (2024)

Artificial Intelligence will profoundly reshape modern society, transforming workplaces, entire industries, and our everyday lives while captivating – and in many cases unnerving – the world with its rapid advancements. Looking forward, AI's trajectory points to continued exponential growth, with many experts anticipating the emergence of Artificial General Intelligence (AGI) in the next years.

This paper aims to forecast the profound impact of AGI, serving as a sequel to my five-part *Foresight* series on forecasting the impact of AI (Makridakis, 2017; 2018a; 2018b; 2018c; 2018d). To illustrate the remarkable capabilities of contemporary AI, I provided a generative Large Language Model (LLM) the following prompt:

‘I am Spyros Makridakis, and I want to publish a sequel to my five-part series on ‘Forecasting the Impact of AI.’ Could you

provide a short summary of the five published articles and suggest an outline for the sequel paper, ‘Forecasting the Impact of AGI?’

The LLM's response, reproduced in the Appendix that follows, was generated in mere seconds. It delivered concise, accurate summaries of each article, demonstrated a deep understanding of their content and interconnections, and proposed a structured outline for the new paper. Accomplishing this otherwise would have required a human expert spending considerable time, effort, and careful thinking.

THE ACCELERATING PATH OF AGI


With the benefit of hindsight, how did my 2017-18 predictions about the impact of AI hold up in February 2026? First, what I got remarkably right is the prediction that exponential progress would continue

From capability to impact

“Frontier models have reached or exceeded human-level expertise on cognitive tasks, **thus** we are on the threshold of AGI, **so** we need to study consequences”

From capability to impact

“Frontier models have reached or exceeded human-level expertise on cognitive tasks, **thus** we are on the threshold of AGI, **so** we need to study consequences”



From capability to impact

“Frontier models have reached or exceeded human-level expertise on cognitive tasks, **thus** we are on the threshold of AGI, **so** we need to study consequences”

Digital tasks to achieve goals



Research



Content creation



Coding



Computer tasks



Benchmark performance

(exams / coding / reasoning / tasks)

From capability to impact

“Frontier models have reached or exceeded human-level expertise on cognitive tasks, **thus** we are on the threshold of AGI, **so** we need to study consequences”

Digital tasks to achieve goals



Research



Content creation



Coding



Computer tasks


Societal impact

(jobs / inequality / GDP / welfare)




Benchmark performance

(exams / coding / reasoning / tasks)




From capability to impact




From capability to impact

The higher we go, the
less clear the
evidence becomes




From capability to impact

Good evidence →




From capability to impact





Good evidence →

Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.0 incorporates 10 evaluations: GDPval-AA, r²-Bench Telecom, Terminal-Bench Hard, SciCode, AA-LCR, AA-Omniscience, IFBench, Humanity's Last Exam, GPQA Diamond, CntPI




From capability to impact




Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.0 incorporates 10 evaluations: GDPval-AA, r²-Bench Telecom, Terminal-Bench Hard, SciCode, AA-LCR, AA-Omniscience, IFBench, Humanity's Last Exam, GPQA Diamond, CnIP






Good evidence →




From capability to impact

How AI assistance impacts coding speed and skill formation



Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.0 incorporates 10 evaluations: GDPval-AA, r²-Bench Telecom, Terminal-Bench Hard, SciCode, AA-LCR, AA-Omniscience, IFBench, Humanity's Last Exam, GPQA Diamond, CMIPI



Good evidence →

Societal impact
(jobs / inequality / GDP / welfare)

Sector / macro productivity


Firm-level performance

Productivity through automation

Usable products


Benchmark performance
(exams / coding / reasoning / tasks)

From capability to impact




Efficiency gains are strong, but revenue impact lags

Globally, AI continues to deliver clear improvements in efficiency; however, the financial returns of this technology are being felt much more slowly. While 74% of organisations surveyed hope AI will grow revenue, only 20% can currently demonstrate measurable impact. This gap shows where many initiatives are stalling. Boards are asking what AI drives beyond just efficiency, and which new offerings and service models it can unlock.




Conflicting evidence →

Good evidence →




From capability to impact

No clear way to measure (yet)




Conflicting evidence

Good evidence




Phase of AI use among organizations using AI in 2025



Efficiency gains are strong, but revenue impact lags


Globally, AI continues to deliver clear improvements in efficiency; however, the financial returns of this technology are being felt much more slowly. While 74% of organisations surveyed hope AI will grow revenue, only 20% can currently demonstrate measurable impact. This gap shows where many initiatives are stalling. Boards are asking what AI drives beyond just efficiency, and which new offerings and service models it can unlock.



Impact needs foundations not FOMO

For true impact, each step needs good foundations.


Otherwise it is hype driven by FOMO (and AI Psychosis).



Impact needs foundations not FOMO

For true impact, each step needs good foundations.

Otherwise it is hype driven by FOMO (and AI Psychosis).



Societal impact

(jobs / inequality / GDP / welfare)

Sector / macro productivity

Firm-level performance

Productivity through automation

Usable products



Benchmark performance

(exams / coding / reasoning / tasks)


Impact needs foundations not FOMO

For true impact, each step needs good foundations.

Otherwise it is hype driven by FOMO (and AI Psychosis).



Three overlapping phases



1. Augmentation Phase (ongoing through the late 2020s): As of early 2026, AI agents (autonomous, goal-oriented systems capable of planning, reasoning, using tools, and executing complex multistep workflows) are already transforming core business operations. Firms across industries are adopting these intelligent digital coworkers at scale to automate high-volume, repetitive tasks. Examples include

2. Substitution Phase (late 2020s to mid-2030s): As AI systems achieve robust, cross-domain reasoning they will surpass humans in speed, accuracy, consistency, and scale. This will first occur across most middle-management functions, and later across upper-management roles. Tasks such as long-term strategic


3. Autonomous Phase (mid-2030s onward): This phase envisions fully autonomous AGI-driven entities including highly evolved Decentralized Autonomous Organizations (DAOs) and fully agentic enterprises, operating with minimal human interventions. These systems would run via AGI algorithms continuously optimized toward predefined objectives, governing themselves through blockchain-

First, augmentation

1. Augmentation Phase

ongoing through the late 2020s

1. Augmentation Phase (ongoing through the late 2020s): As of early 2026, AI agents (autonomous, goal-oriented systems capable of planning, reasoning, using tools, and executing complex multistep workflows) are already transforming core business operations. Firms across industries are adopting these intelligent digital coworkers at scale to automate high-volume, repetitive tasks. Examples include




Then: substitution due to better models



1. Augmentation Phase

ongoing through the late 2020s


Better models!




Then: substitution due to better models




Then: substitution due to better models




The Assumption:
Models better/faster at reasoning,
thus they replace decision makers.



Then: substitution due to better models



The Assumption:
Models better/faster at reasoning,
thus they replace decision makers.




"I think that if you work as a radiologist, you are like Wile E. Coyote in the cartoon. You're already over the edge of the cliff, but you haven't yet looked down. There's no ground underneath. People should stop training radiologists now. It's just completely obvious that in five years deep learning is going to do better than radiologists."

Nov 24, 2016


As of May 2026, the demand for radiologists is at an all-time high, with job growth expected to rise 9% through 2034—significantly faster than the average for all occupations. Contrary to predictions a decade ago that AI would make radiologists obsolete, AI has instead become an essential tool that "supercharges" radiologists by handling high-volume tasks, leading to increased productivity and higher salaries, which reached an average of \$571,000 in 2025. Fortune +2

Then: substitution due to better models




The Assumption:
Models better/faster at reasoning,
thus they replace decision makers.



On April 8, 2026, Sam Altman confirmed that Codex CLI had reached 3 million weekly active users. The growth trajectory since the beginning of 2026 has been steep:




GPT-5.5, released in April 2026, sets a high-water mark for AI, achieving expert-level reasoning that outperforms previous models and **often rivals senior-level human performance in coding and technical tasks**. It excels in speed,, reducing 12 hours of human expert work to just 10 minutes, changing the cost dynamics of AI usage. [MindStudio +4](#)



And finally: full autonomy



And finally: full autonomy




Huge uncertainty for such predictions.

- 10 years in AI years is almost a human century
 - New best model every 2 months
- Current models work best on verifiable tasks
 - Eg best models are coding models
- Architectural breakthroughs are needed
 - Solve continuous learning
 - Solve long term planning
 - Solve memory
- What fuels these autonomous entities?


A more “naive forecast”

Software engineering job openings soar



A more “naive forecast”

Software engineering job openings soar





Software gets ~x% cheaper

- Requirements gathering
- Prototyping
- More productive engineers

A more “naive forecast”

Software engineering job openings soar




Software gets ~x% cheaper

- Requirements gathering
- Prototyping
- More productive engineers

With Great Substitution, Comes Great Augmentation

For every job at-risk of substitution, there are other jobs that stand to benefit

High Exposure to AI Substitution



High Exposure to AI Augmentation



Source: Goldman Sachs Global Investment Research (April 2026)

<https://a16z.com/the-ai-job-apocalypse-is-a-complete-fantasy/>

Bonus: The real problem with the labour market.





Figure 7: New job starts among workers age 22-25 in occupations with high observed exposure and no AI exposure, Current Population Survey

The top panel shows the percent of young workers starting new jobs in high vs. no exposure occupations. The bottom panel measures the gap between these two series in a difference-in-differences framework.


Bonus: The real problem with the labour market.




Question:
What does this mean
for education?

The expectation loop & the economics of AI

The expectation loop & the economics of AI



The expectation loop & the economics of AI



The expectation loop & the economics of AI



The expectation loop & the economics of AI



Some problems with the expectation loop

Raises at 50B\$



Bloomberg.com

<https://www.bloomberg.com/news/articles/ai-codin...>

AI Coding Startup Cursor Plans New Model to Rival ...



Mar 19, 2026 — **Cursor** has also been in talks to **raise** a new round of financing at a roughly **\$50 billion** valuation, **Bloomberg** News reported this month. [Read more](#)



Some problems with the expectation loop

Raises at 50B\$



Bloomberg.com

<https://www.bloomberg.com/news/articles/ai-codin...>

AI Coding Startup Cursor Plans New Model to Rival ...

Mar 19, 2026 — **Cursor** has also been in talks to **raise** a new round of financing at a roughly **\$50 billion** valuation, **Bloomberg** News reported this month. [Read more](#)



Composer 2 includes our first continued pretraining run and scaled reinforcement learning

Model	Terminal-Bench 2.0 score
GPT-5.4	75.1%
Composer 2	61.7
Opus 4.6	58.0
Opus 4.5	52.1
Composer 1.5	47.9

Trains their in-house model to rival big labs

Some problems with the expectation loop

Raises at 50B\$



Bloomberg.com

<https://www.bloomberg.com/news/articles/ai-codin...>

AI Coding Startup Cursor Plans New Model to Rival ...

Mar 19, 2026 — **Cursor** has also been in talks to **raise** a new round of financing at a roughly **\$50 billion** valuation, **Bloomberg** News reported this month. [Read more](#)



Composer 2 includes our first continued pretraining run and scaled reinforcement learning

Model	Terminal-Bench 2.0 score
GPT-5.4	75.1%
Composer 2	61.7
Opus 4.6	58.0
Opus 4.5	52.1
Composer 1.5	47.9

Trains their in-house model to rival big labs



Kimi.ai @Kimi_Moonshot · 14 ü


Congrats to the [@cursor_ai](#) team on the launch of Composer 2!

We are proud to see Kimi-k2.5 provide the foundation. Seeing our model integrated effectively through Cursor's continued pretraining & high-compute RL training is the open model ecosystem we love to support.

Note: Cursor accesses Kimi-k2.5 via [@FireworksAI_HQ](#) ' hosted RL and inference platform as part of an authorized commercial partnership.

Its fine-tuned Kimi k2.5 Moonshot is valued at ~4.3B

Where the problems begin.



Where the problems begin.

What about the
problems here



Physical world



Human goals:

- Improve quality of life
- Earn money
- Publish more
- Other...


Instructions /
Feedback



Actions /
Clarification



Digital world




Digital tasks to achieve goals

-  Research
-  Content creation
-  Coding
-  Computer tasks

AGI
happens
here

Where the problems begin.



Where the problems begin.


Physical world



Human goals:


- Improve quality of life
- Earn money
- Publish more
- Other...

Instructions /
Feedback




Actions /
Clarification

Digital world



Digital tasks to achieve goals

-  Research
-  Content creation
-  Coding
-  Computer tasks



Where the problems begin.

Physical world



6. Chips / data centers / talent



Where the problems begin.

Physical world



6. Chips / data centers / talent



 Reuters
<https://www.reuters.com/markets/commodities/us-...>

[US AI boom faces electric shock](#)

Feb 25, 2026 — Big Tech's race to dominate artificial intelligence **may soon hit a nasty road bump** as U.S. electricity grids struggle to keep pace with the ...

 Nature
<https://www.nature.com/news>

[Boycott of major AI conference exposes a growing US– ...](#)

Apr 14, 2026 — Boycott of major AI conference exposes a growing US–China divide. **China-based researchers make up more than half of all lead authors at NeurIPS.** [Read more](#)

 WSJ
<https://www.wsj.com/tech/spacex-google-in-talks-to-...>

[SpaceX and Google Are in Talks to Launch Data Centers ...](#)

18 hours ago — A launch deal would put the two companies in partnership as they gear up to compete on orbital **data centers**, an unproven technology that **SpaceX** ... [Read more](#)


 Financial Times
<https://www.ft.com/.../Asia-Pacific/economy/China>

[Nvidia AI chip sales to China stalled by US security review](#)

Feb 3, 2026 — **Nvidia's sales of H200 AI chips to China are still awaiting final approval** from Washington nearly two months after Donald Trump greenlit ... [Read more](#)

Where the problems begin.

Question:
What if this part
slows down?



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

The AI supply chain




Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

The AI supply chain



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

The AI supply chain

Chips



AMD



ASML

Cloud



Models



OpenAI

ANTHROPIC



Google DeepMind

Meta



7

Applications /
agents

Built on models from



OpenAI

ANTHROPIC

Google DeepMind

Meta



AI Assistant

ChatGPT

Gemini



Code Agent

Claude Code

GitHub Copilot



Data Analyst Agent

ChatGPT Data Analyst

Gemini

8

Firms / workers /
society



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

The AI supply chain

Hypothetical questions:


Q1: What is the % of these companies in the stock market?

Q2: What happens to the stock market without these companies?

Q3: What about potential disruptions in the supply chain?

7

Applications /
agents



8

Firms / workers /
society



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?


Market Summary > NVIDIA Corp

118.58 USD

-20.64 (-14.83%) 4 past 5 days

Closed: 27 Jan 5:07 PM GMT-5+ Disclaimer
After hours 120.66 +2.08 (1.76%)

1D 5D 1M 6M YTD 1Y 5Y Max



Follow

deepseek

DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI
research@deepseek.com
Abstract

The AI supply chain



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion of Ethiopia, the second most African nation. What makes this even more intriguing is that Nvidia's 2025 revenues were generated by 36,000 employees, resulting in a contribution of \$3,600 per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker.

Reuters
https://www.reuters.com › business › energy › white-ho...
White House scrambles for gas-price relief as Iran war ...
3 hours ago — ... prices even higher this summer. Energy prices have spiked since Iran cut off access to the Strait of Hormuz, a waterway that normally ... [Read more](#)

Politico
https://www.politico.com › news › 2026/05/14 › taiwan...
Taiwan emerges as flash point in Trump-Xi talks
23 hours ago — BEIJING — The sensitive issue of Taiwan intruded Thursday on a high-stakes meeting between U.S. President Donald Trump and Chinese leader Xi ... [Read more](#)

Reuters
https://www.reuters.com › business › media-telecom › c...
China's Huawei hypes up chip and computing power plans ...
Sep 18, 2025 — The announcement broke years of secrecy about its chips business and could increase the stakes in the U.S.-Sino war for tech supremacy.

Market Summary > NVIDIA Corp


118.58 USD

-20.64 (-14.83%) 4 past 5 days

Closed: 27 Jan 5:07 PM GMT-5+ Disclaimer

After hours 120.66 +2.08 (1.76%)

1D 5D 1M 6M YTD 1Y 5Y Max



[Follow](#)

[deepseek](#)

DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI
research@deepseek.com

Abstract

The AI supply chain

1 Energy / grid / cooling / land



2 Semiconductors / fabs / packaging



3 AI chips + networking



nvidia
GPU + networking + CUDA + AI factory design


4 CUDA / software ecosystem




5 Cloud / data centers



6 Models



7 Applications / agents



8 Firms / workers / society



Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

Question 1: Who owns the technology and who gets access?

Honestly at this point, the bottleneck does not feel to be something technical.

I am the bottleneck. Me ,my cognitive limitations, and my salary that I can't afford 5 Claude max subscriptions.

I can empty my to-do lists in a night, I can build anything I want.

I am only limited by the quality of the high level abstract ideas I generate.

7:52 μ.μ. ✓✓

Ownership and inequality

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

Question 1: Who owns the technology and who gets access?

Once an AI system exceeds human-level intelligence, if granted the capability of self-improvement it could set its own goals that may conflict with those of humans. This creates an *alignment problem*: the challenge of designing a system with objectives that remain stable and human-compatible even as the system undergoes radical self-modification and acquires power far beyond that of human capabilities.

Question 2: Who aligns the aligners?



BBC

<https://www.bbc.com/news/articles>

Judge rejects Pentagon's attempt to 'cripple' Anthropic

Mar 26, 2026 — A federal judge told the government it could not immediately enforce a ban on Anthropic's tools.



OpenAI

<https://openai.com/index/our-agreement-with-the-de...>

Our agreement with the Department of War

Mar 2, 2026 — Yesterday we reached an agreement with the Pentagon for deploying advanced AI systems in classified environments, which we requested they also ... [Read more](#)



Anthropic

<https://www.anthropic.com/news/where-stand-depar...>

Where things stand with the Department of War

Mar 5, 2026 — Yesterday (March 4) Anthropic received a letter from the Department of War confirming that we have been designated as a supply chain risk to ... [Read more](#)

Honestly at this point, the bottleneck does not feel to be something technical.

I am the bottleneck. Me ,my cognitive limitations, and my salary that I can't afford 5 Claude max subscriptions.


I can empty my to-do lists in a night, I can build anything I want.

I am only limited by the quality of the high level abstract ideas I generate.

7:52 p.m. ✓

Wrapping up. Why forecasting “AGI” is hard.

1. Trends lean towards augmentation -> Cheap software + more digitalization
2. No striking productivity gains (firm level, not worker level). These will show during consolidation and mass augmentation.
3. Physical boundaries should be considered.
4. No evidence of a significant job market shift (yet)
5. Revenue and valuation might be weak predictors.



Wrapping up. Why forecasting “AGI” is hard.

1. Trends lean towards augmentation -> Cheap software + more digitalization
2. No striking productivity gains (firm level, not worker level). These will show during consolidation and mass augmentation.
3. Physical boundaries should be considered.
4. No evidence of a significant job market shift (yet)
5. Revenue and valuation might be weak predictors.

Plus: It might be a choice and not a forecast:

“The question is not what AI will do to us, but what we are going to do with AI”, Erik Brynjolfsson, Stanford economist.




From student anxieties to vice coding, Brynjolfsson frames A.I. as a tool that amplifies intention rather than replaces workers. Photo by Ray Rostov/Gallo Images for Unsplash/Lia.

<https://observer.com/2026/04/stanford-erik-brynjolfsson-ai-labor/>

What that means for forecasting


What that means for forecasting

- The practise of forecasting can be explained by the same framework.




What that means for forecasting

- The practise of forecasting can be explained by the same framework.



What that means for forecasting

- The practise of forecasting can be explained by the same framework.
- We just augment our toolbox.
- More and cheaper software
- Who owns the major tools?
- Important to keep up (as much as possible) with new tools



Thank you for the attention 😊



filotas.theodosiou@vives.be



<https://github.com/FilTheo>



<https://www.linkedin.com/in/filotas-theodosiou/>



<https://filtheo.github.io/>

A more “naive forecast”

1. Augmentation Phase

+

2. Substitution Phase

+

3. Autonomous Phase


=

Labour market redistribution

The asymmetric impact of generative AI

Researchers observed a drop in job postings for occupations highly exposed to automation following the launch of ChatGPT. At the same time, they noted a rise in employer demand for jobs likely to be **augmented** by AI.

Average effect (number of job postings)



Source: "Displacement or Complementarity? The Labor Market Impact of Generative AI," HBS Working Paper, August 2025.



Harvard
Business
School

Working
Knowledge

Ownership, inequality &Alignment

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

Ownership, inequality &Alignment

IMPLICATIONS FOR SOCIETIES, FIRMS, AND MANAGERS

Nvidia's 2025 revenues of \$130 billion are not far off from the \$126 billion GDP of Ethiopia, the second most populous African nation. What makes this comparison intriguing is that Nvidia's revenues were generated by 36,000 employees, resulting in a contribution of \$3.6 million per employee; Ethiopia's 2025 GDP was produced by a labor force of 55 million, resulting in \$2,300 per worker. This mind-

Technology?
Or ownership (and influence)?

■ Once an AI system exceeds human-level intelligence, if granted the capability of self-improvement it could set its own goals that may conflict with those of humans. This creates an alignment problem: the challenge of designing a system with objectives that remain stable and human-compatible even as the system undergoes radical self-modification and acquires power far beyond that of human capabilities.

More importantly, who aligns the aligners?

BBC
<https://www.bbc.com/news/articles>

Trump orders government to stop using Anthropic in battle ...

Feb 27, 2026 — Anthropic's tools will be phased out of all government work over the next six months, Trump said. As for Anthropic's other customers, the ... [Read more](#)



Anthropic

<https://www.anthropic.com/news/statement-departm...>

Statement from Dario Amodei on our discussions with the ...

Feb 26, 2026 — I believe deeply in the existential importance of using AI to defend the United States and other democracies, and to defeat our autocratic ... [Read more](#)



OpenAI

<https://openai.com/index/our-agreement-with-the-de...>

Our agreement with the Department of War 🇺🇸

Mar 2, 2026 — Yesterday we reached an agreement with the Pentagon for deploying advanced AI systems in classified environments, which we requested they also ... [Read more](#)